

Abstract

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Acknowledgement

I would like to say "thanks brah"

Contents

1	Introduction	4
1.1	Introduction to game generation	4
1.2	Research Objectives	4
1.3	How to read this thesis	5
1.3.1	Terms	5
2	Existing research and framework	7
2.1	Content generation	7
2.1.1	Game-level generation	7
2.1.2	Other uses of PCG	7
2.1.3	Advantages and disadvantages using PCG	7
2.2	Game generation	8
2.2.1	Something to note: Levels	8
2.3	Framework: VGDL and the GVG-AI framework	9
2.3.1	VGDL game description	9
2.3.2	The GVG-AI framework: AI controllers and testing	10
3	Extending VGDL	13
3.1	Writing new VGDL games	13
3.1.1	Describing existing games in VGDL	13
3.1.2	Results	13
3.2	Creating new VGDL controllers	14
3.2.1	Action-arcade controllers	14
3.2.2	Puzzle controllers	15
3.3	FastVGDL	15
4	Automatic generation of VGDL descriptions	16
4.1	Generating processes	16
4.1.1	Mutation of example games	16
4.1.2	Random game generation	17
4.2	Initial test: Experimental setup	18
4.3	Initial test: Results	18
4.4	Discussion of initial test results	20
5	Evaluation of games (fitness functions)	23
5.1	Fitness function features: Introduction	23
5.2	Fitness function features: Selection	23
5.2.1	Analysis of feature choices	24
5.2.2	Selection and evolution of weights	24
5.2.3	Outcome	24
5.3	Evolution of games: Setup	25
5.4	Evolution of games: Results	25
5.5	Discussion of results	25
5.6	HUUUUSK Designed action-arcade games: Generated levels	25

6	Puzzle generation	26
6.1	Turn-based puzzle games	26
6.1.1	Games	26
6.1.2	Puzzle solving AIs	26
6.1.3	Level generation	26
6.1.4	Level generation for generated games	27
7	Conclusion	28
	Appendices	31
A	GVG-AI example games	32
B	Games designed during thesis	34
B.1	Action arcade games	34
B.2	Puzzle games	34
C	Designed games results	36

Chapter 1

Introduction

Procedural generation of game content (PCG-G) has become a widely used tool for video-game developers. Automatically generating content allow game developers to provide more content, with more differentiation, from a limited supply of assets, most often by helping to produce levels, maps or dungeons for their games. Additionally PCG is often used to maintain a challenge for the player by presenting an unpredictable and renewed game experience on re-playing the games in which it is used.

Rogue [1980] is one of the earliest examples on procedurally generating game levels (or dungeons). Several highly popular games has since used procedural level generators, some as a main part of their game-play (The Seven Cities of Gold [?], Civilization II(-V) [?], Minecraft [2011], Spelunky [2009] and many others), while other has used PCG tools to create finished levels (Rescue on Fractalus! [?], Farcry 2 [?]. PCG has also been used to create a wider array of weapons and items (Borderlands-series [?], Diablo-series [?], Torchlight-series [?], Dark Age of Camelot [?]), or even to create graphical or audio assets (RoboBlitz [?], .kkrieger [?]).

1.1 Introduction to game generation

An interesting next step to take is to not only generate content for a game, but to generate completely new games and game-play – i.e. to both generate a set of rules and mechanics, and generate the content necessary to play the game (levels, sprites etc.).

A possible way of automatically generating complete games might be to search through a space of games represented in a programming language like C or Java. However, the proportion of programs designed in such languages that can actually be considered a game is rather small (and much less, the amount of human-enjoyable games). Searching through programs defined in a game description language (GDL), designed to encode games (and only games), can increase the density of well-defined games.

Even searching a fairly well defined space of games still expect that we have some way of automatically telling good games from bad games (or at least, not-very-good games from really bad games). In other words, we need a fitness function. A fitness function could partly consist of inspecting the generated rules as expressed in the GDL, e.g. to make sure that the player can interact with the game in some way, and that there are winning- and losing conditions, which could in principle be fulfilled. However there are many bad games that fulfil such criteria.

To really understand a game it seems you need to play it. The fitness function therefore needs to incorporate a capacity to automatically play the games it is evaluating.

1.2 Research Objectives

The main objectives of this study was to construct a generation process able to create simple games and levels using the game description language (VGDL), with the goal of being as enjoyable for human players as possible. The focus of the study was to create a fitness function able to differentiate between games (and levels) of different quality, to be able to search and sort a set of generated games, and thereby create interesting content.

The main approach used to create the fitness function was to use the results of a series of general game-playing (knowledge free) algorithms, and using the collection of their performance profiles to construct individual fitness features. A series of human-designed games, assumed to be human-enjoyable, was used as a baseline for "good games" and was intensely used in choosing and weighing fitness features. This approach was focused on a specific genre and type of games (*arcade-action games*), which is in the focus of the framework used (the GVG-AI framework). Additionally this type of games fit very well with general game-playing algorithms used, which are not very good at exploring deep into a space of game-states (which is often not needed in the game-genre – quick movements to avoid enemies and increase the score is in the focus).

The second part of this thesis is focused on games of the *puzzle*-genre described in VGD, which overall has a much smaller game-state space and so potentially are more well-suited for a more in-depth analysis. This approach is mainly focused on generating new game-levels for either existing games, or generated VGD programs, by completely searching through all possible states of a game (using a breadth-first algorithm), finding all possible solutions, with the goal of being able to judge the difficulty, and thereby the quality of a game/level pair.

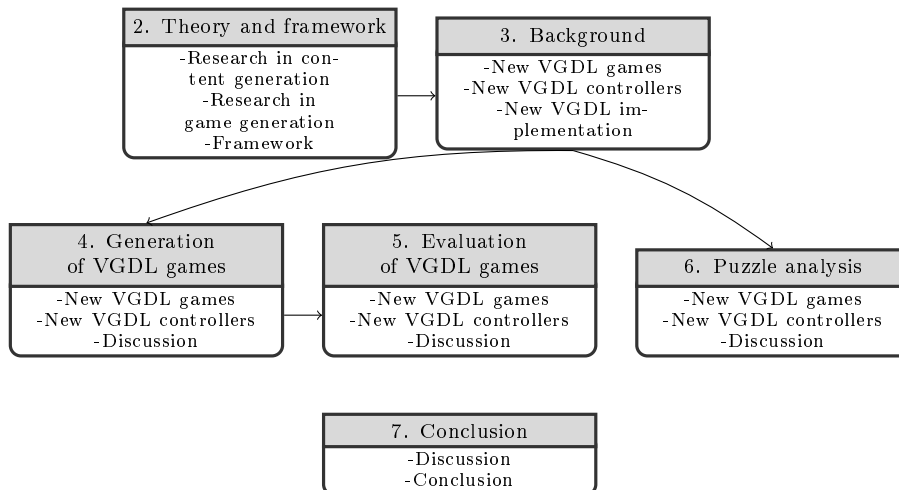
Assumptions

As mentioned above, we assume that a subset of the human-designed VGD games described (Section 2.3, 3.1) – the ones which are interpretations of commercial games – are of high quality (i.e. enjoyable for human players). This is not always clear when actually playing the games in the framework used, since they mostly need crucial elements like, audio, graphic effects and a proper control-scheme fitting to the game.

We additionally assume that completely randomly generated games using the VGD language (without sorting/searching using a fitness function) produce games with an extremely (almost negligible) low chance of being human-enjoyable. This assumption is mostly made from initial tests into game generation using the VGD language.

1.3 How to read this thesis

In the following chapters we will define the problem in more detail and explain the existing research and framework used. We will then show the results of the seeking to fulfil the different research goals, and finally discuss the conclusion of the results. The figure below shows the structure of the thesis, explaining the reasoning behind each chapter and how it fits in to the larger picture.



1.3.1 Terms

This project is focused on games of two restricted genres and types:

1) Single-player, 2-dimensional, top-down, action games, in which the player controls a single avatar which can (maximally) be moved in the four cardinal directions, possibly with an "action"- (shoot etc.) button. These games will be referred to as **Action-arcade games**.

And: 2) Single-player, 2-dimensional, top-down, turn-based puzzle games, in which the player controls a single avatar which can (maximally) be moved in the four cardinal directions, possibly with an "action"-(shoot etc.) button., and all events occur as a result of player actions (i.e. no NPC's or falling boulders). This type of games will be referred to as **Puzzle games** in the project.

Chapter 2

Existing research and framework

In this chapter we describe the existing research on the topic of procedural content generation (PCG) for games, and specifically game- and level-generation, which is of relevance to the project. Additionally the framework used for game- and level-generation, and testing is described thoroughly.

2.1 Content generation

Before discussing generation of complete games, the research on procedural content generation (PCG) for games are investigated. The task of PCG is to create elements for specific parts of a game, for instance levels, weapons, environments, textures and sounds. Hendrikx et al. [2013] proposes a layered taxonomy to describe PCG of different depths in games, while performing a survey of different ways in which PCG has been implemented in commercial games, noticing that some layers has not been thoroughly explored.

Several different approaches are used in PCG (both in commercial games and in research), however the "search-based" approach (explored by Togelius et al. [2011], Nelson et al. [2014a]) dominates the problem. In search-based generation a *fitness function* is used to score generated content by their quality, making it possible to ensure a high quality of the generator, by either simply generating a large amount of content and choosing the best, or using an evolutionary strategy to evolve better content.

2.1.1 Game-level generation

Almost all video games are highly dependent on designed level or map (or levels/map) in which the game play takes place. Levels are often responsible for creating increasingly challenging situations, and thereby entertaining game play. Super Mario would be way less interesting if enemies and platforms were spread randomly across each level, which would almost certainly result in the game being either too easy or too difficult.

Several scientific projects have focused on the topic of level-generation with an array of different approaches, some requiring player-input, or simply a helper for human level-designer.

In addition to the other game genres mentioned above, a lot of work has been done in generating levels for puzzle games (ie. Sokoban and Cut The Rope). Most of the work is focused on generating levels for single game, where some attempt to create general level generation procedures.

2.1.2 Other uses of PCG

Other games have been published applying PCG in different ways. In *Borderlands 2* [2012] the player have access to an extreme amount of different weapons, designed by automatic generation.

2.1.3 Advantages and disadvantages using PCG

Having a procedure for generating content in games has a number of advantages:

- Content generation allows games to potentially have an endless amount of possibilities, often resulting in higher re-play value than linear games.

- PCG can help game designers generate game levels, often simply by letting the designer explore possible layouts, add decorations or finishing up human-designed levels.

PCG also has some problems for generating content for levels:

- Generated content can have a feeling of being unauthentic and/or too random. For instance, the positions of decorations in a level (barrels, paintings etc.) can be important for humans, while they are not for computers.
- In games where the player follows a story it can be difficult to automatically generate levels, that follow what the game designer intend to happen

2.2 Game generation

Generating complete games through algorithms is a problem that is being increasingly researched, but work has been done on the topic for last decade. Because the problem is in general quite large, a subset of the problem is usually handled: Only generating certain types of games-, or using different (restricted) frameworks. Video games may consist of a large number of tangible and intangible components, including rules, graphical assets, genre conventions, cultural context, controllers, character design, story and dialog, screen-based information displays, and so on Cook and Colton [2014], Liapis et al. [2014], Nelson and Mateas [2007].

In this project we look specifically at generating the game-play and -setting of games, i.e. defining a set of game-rules and -objects, and specifying the levels in which the game-play takes place. The two main approaches that have been explored in generating game rules are reasoning through constraint solving Smith and Mateas [2010] or search through evolutionary computation or similar forms of stochastic optimisation Togelius and Schmidhuber [2008], Browne [2008], Font et al. [2013]. In either case, rule generation can be seen as a particular kind of procedural content generation Nelson et al. [2014b].

It is clear that generating a set of rules that makes for an interesting and fun game is a hard task. The arguably most successful attempt so far, Browne's Ludi system, managed to produce a new board game of sufficient quality to be sold as a boxed product Browne [2008]. However, it succeeded partly due to restricting its generation domain to only the rules of a rather tightly constrained space of board games. A key stumbling block for search-based approaches to game generation is the fitness/evaluation function. This function takes a complete game as input and outputs an estimate of its quality. Ludi uses a mixture of several measures based on automatic playthrough of games, including balance, drawishness and outcome uncertainty. These measures are well-chosen for two-player board games, but might not transfer that well to video games or single-player games, which have in a separate analysis been deemed to be good targets for game generation Togelius et al. [2014]. Other researchers have attempted evaluation functions based on the learnability of the game by an algorithm Togelius and Schmidhuber [2008] or an earlier and more primitive version of the characteristic that is explored in this paper, performance profile of a set of algorithms Font et al. [2013].

A typical approach for generating complete games is searching in a space of possible games. This basically requires two things: That the ratio of enjoyable games in the set is not too low - otherwise those games might never be found. To increase this ratio a game description language (GDL) is often used (searching through all possible Java or C programs would lead to an enormous amount invalid games). Also, to search through a set of games it is necessary to be able to calculate a fitness value for each game, valuing how enjoyable the game is.

2.2.1 Something to note: Levels

The problem of generating complete games is heavily linked to the problem of generating levels since most games are deeply tied to the geometry and object-placement defined in levels. A part of the success of Ludi stem from the fact that levels (boards) were as much in focus, as the game-rules themselves.

2.3 Framework: VGDL and the GVG-AI framework

Regardless of which approach to game generation is chosen, one needs a way to represent the games that are being created.¹ For a sufficiently general description of games, it stands to reason that the games are represented in a reasonably generic language, where every syntactically valid game description can be loaded into a specialised game engine and executed. There have been several attempts to design such GDLs. One of the more well-known is the Stanford GDL, which is used for the General Game Playing Competition Genesereth et al. [2005]. That language is tailored to describing board games and similar discrete, turn-based games; it is also arguably too verbose and low-level to support search-based game generation. Another attempt at an VGDL is called PuzzleScript, created by game designer Stephen Lavelle. The language (as its name suggest) is focused on turn-based puzzle games, but the engine allows for simple forms of animation and movement. The language is relatively high level compared to Stanford GDL, but does not support that large set of video games.

2.3.1 VGDL game description

The various game generation attempts discussed above feature their own GDLs of different levels of sophistication; however, there has not until recently been a GDL for suitably for a larger space of video game types- and genres. The Video Game Description Language (VGDL) is a GDL designed to express 2D arcade-style video games of the type common on hardware such as the Atari 2600 and Commodore 64. It can express a large variety of games in which the player controls a single moving avatar (player character) and where the rules primarily define what happens when objects interact with each other in a two-dimensional space. VGDL was designed by a set of researchers Levine et al. [2013], Ebner et al. [2013] (and implemented by Schaul Schaul [2013]) in order to support both general video game playing and video game generation. In contrast to other GDLs, the language has an internal set of classes, properties and types that each object can defined by, which the authors suggest the user to extend.

Objects have physical properties (i.e. position, direction) which can be altered either by the properties defined, or by interactions defined between specific objects. Playing the games also required a specified level which defines a set of game-tiles deciding the initial locations of sprites. When running games sprite can move in between game-tiles (if the **speed** is a fraction).

A VGDL description has four parts:

SpriteSet Defines which sprites can appear in the game. Each sprite must be designated by a class, in which a set of predefined actions exists. Also a set of parameters can be fed to each sprite, configuring for instance the speed or how often the sprite takes an action, and the parameters used for the different sprite-classes – for instance, for the class **Flicker** (a simple extension to the base sprite, the sprite is destroyed after a specified amount of time) the lifetime can be adjusted. Sprites can additionally be designed in a tree structure, where multiple sprites have the same parent sprite, making more possibility for the interactions and terminations described below.

InteractionSet Each line defines what happens when a set of two sprites collide with each other (located on the same game-tile). Each interaction is represented by a class defining the action to take (i.e. push back, or kill sprite), and a set of parameters specific to each interaction-class, as well as the score achieved for letting the interaction happen.

TerminationSet Defines how the game can end. Each line in this set has a win parameter, which is set to true or false; winning or losing the game. Each termination-function is represented by a class defining under which conditions the game should end, which can for instance when there exists 0 of a certain sprite (when all of the sprites are killed).

LevelMapping The job of the **LevelMapping** is to translate from a character (**char**) in a level-file (explained below), to sprites from the **SpriteSet**. A single mapping can be shared by several sprites, causing the sprites to be created on the same game-tile.

Additionally each game can only be played using level files. A level file is written using the **LevelMapping** characters, where each character defines a tile of the game, and spaces defines empty tiles.

¹See Nelson et al. [2014b] for a discussion of game-rule representation choices.

```

1 BasicGame
2   SpriteSet
3     city > Immovable color=GREEN img=city
4     explosion > Flicker limit=5 img=explosion
5     movable >
6       avatar > ShootAvatar stype=explosion
7       incoming >
8         incoming_slow > Chaser stype=city color=ORANGE speed=0.1
9         incoming_fast > Chaser stype=city color=YELLOW speed=0.3
10
11   LevelMapping
12     c > city
13     m > incoming_slow
14     f > incoming_fast
15
16   InteractionSet
17     movable wall > stepBack
18     incoming city > killSprite
19     city incoming > killSprite scoreChange=-1
20     incoming explosion > killSprite scoreChange=2
21
22   TerminationSet
23     SpriteCounter stype=city win=False
24     SpriteCounter stype=incoming win=True

```

Figure 2.1: Example of VGDL description - a simple implementation of the game Missile Command

```

1 w   m   m   m   m   m   m   nw
2 w           w
3 w           w
4 w           w
5 w           w
6 w           w
7 w           A   w
8 w           w
9 w           w
10 w          w
11 w   c   c   c   c   c   c   c   w
12 wwwwwwwwwwwwwwwwwwwwwwwwwwwww

```

Figure 2.2: Example of VGDL level description - a level for the implementation of the game Missile Command

2.3.2 The GVG-AI framework: AI controllers and testing

The GVG-AI framework is a testbed for testing general game-playing controllers against games specified using VGDL. Controllers are called once at the beginning of each game for setup, and then once per clock tick to select an action. Controllers do not have access to the VGDL descriptions of the games. They receive only the game’s current state, passed as a parameter when the controller is asked for a move. However these states can be forward-simulated to future states. Thus the game rules are not directly available, but a simulatable model of the game can be used.

GVG-AI sample controllers

The GVG-AI contains a series of sample AI controllers, which some were used in testing games in the project. Below is a short description of each of the sample controllers used:

MCTS “Vanilla” MCTS using UCT.

GA Uses a genetic algorithm to evolve a sequence of actions.

OneStep-Heuristic Heuristically evaluates the states reachable through one-step lookahead. The heuristic takes into account the locations of NPCs and certain other objects.

Random Chooses a random action from those available in the current state.

GVG-AI example games

The framework additionally contains 20 hand-designed games, which partly consist of interpretations of classic video games (e.g. Boulderdash, Frogger, Missile Command and Pacman), while some are original creations by the General Video-Game AI Competition’s organizers. Note however, that the interpretations of class games all have different stages of simplifications to them, causing several of the games’ features to be missing – for instance, the player cannot push boulders in the VGDL version of Boulderdash and they do not "roll" as in the original, and the ghosts in Pacman behave in a much more simplified fashion. A short summary of each of the example games can be seen in Appendix A

The games can be described (except for the game *Sokoban*) as *action-arcade* games, in that the player controls a single avatar which must be moved quickly around in a 2D-setting to win, or to get a high score (the player is able to increment a score counter in all of the games). Figure 2.1 shows the VGDL description of the game *Missile Command*, and 2.2 a level description from the same game.



Figure 2.3: A visual representation of a few of the VGDL example games. From top-left: *Zelda*, *Portals* and *Boulderdash*

Restrictions of VGDL and the GVG-AI framework

The VGDL implementation of the GVG-AI competition is essentially (without great extension) only able to describe certain action-arcade games, and puzzle games. For instance, the lack of possibility to traverse different levels make **Adventure** games impossible to make, and the restriction of only being able move a single character with the keyboard (the avatar) makes great restriction in creating **Strategy** games.

In this project we will focus on the two main type of games describable in VGDL: Action-arcade games and turn-based puzzle games. We make the simple distinction that arcade games contain elements (sprites) that move by themselves, whereas interactions only occur as a result of the

avatar moving in puzzle games. This puts the game *Portals* in the arcade genre, even though it contains puzzle elements.

Chapter 3

Extending VGDL

This chapter describes a series of extensions we constructed for VGDL and the GVG-AI framework, making a more in-depth analysis possible. We created a series of VGDL games, AI controllers and implemented a new, simplified framework to play through *puzzle games* with less time and memory use.

3.1 Writing new VGDL games

To increase the size of the set of designed games, to allow for a more precise analysis of what makes a game "good", we created fourteen new game descriptions in VGDL.

Another important reason for creating new descriptions, was to introduce a series of *puzzle games* to be analysed, since, as mentioned in Section 2.3.1 the example games from the GVG-AI competition are almost all *action-arcade* games.

3.1.1 Describing existing games in VGDL

The main goal when developing a game - which also applies in this case - is to ensure that it is enjoyable to human-players. Therefore the games implemented are all interpretations of published existing games (and not original creations). As described in **Section 2.3.2** VGDL can only describe relatively simple games of certain types/genres, and so only a limited number of games can be translated without severely changing the game-play.


The games we created, and described below are in general "exact" copies of the originals - containing all the game-play features and interactions but lacking elements like audio, graphics and controls. However many lack certain (non-essential) game-play features like bonus points (like fruits in Pac-Man), infrequently appearing enemies (UFOs in Space Invaders) or features which only appear in some levels of the games.

3.1.2 Results

A series of fourteen commercial games was re-created in VGDL to be used in future tests. Figure 3.1 shows one of the games translated to VGDL. Minor changes to the GVG-AI framework was made for a small number of the games, to be able to correctly copy the games' core gameplay features. The original game levels were additionally translated into VGDL level description, with five levels being created for each game. For some of games the levels could not be completely copied, for instance because the levels were too large in size (*Bolo Adventures*), or because a lack of game features implemented in the clones (*Chip's Challenge*), and the resulting levels are therefore simplifications of the originals..

Action-arcade games

A set of four Atari -arcade and -2600 games were found to be suitable for interpretation; *Solar Fox*, *Crackpots*, *Centipede* and *Astrosmash*, and a VGDL game description was written for each. A description of each of the games and their interpretations can be seen in Appendix B

Figure 3.1: Interpretation of the classic arcade game *Centipede* [1983]

Puzzle games

A set of puzzle games from several different platform, all featuring a single avatar, was found to be suitable to be described in VGDL. Several of these games can be described as *Sokoban*-clones with different spins on the gameplay. The games interpreted were *Sokoban*, *Bait*, *Bombuzal*, *Bolo Adventures*, *Zen Puzzle*, *The Citadel*, *Brainman*, *Chip's Challenge*, *Modality* and *Painter*.

3.2 Creating new VGDL controllers

To be able to probe VGDL games in more detail, we created a series of new AI controllers with various approaches to finding a results. In total five new controllers was generated: Three to play the action-arcade style games, and two only focused on puzzle games.

3.2.1 Action-arcade controllers

Three controllers of an increasing degree of cleverness was introduced: *One-step* (least clever), *Deep-search* and *Explorer* (most clever). The controllers greatly differ in their strategy in simulating the forward model (accessible when running the GVG-AI framework) using different actions, and different series of actions.

One-step

The One-step AI only attempt to advance the forward model once, for each allowed action in the game. The score and win/lose state of the resulting game states are then considered, and the action with the resulting best state is chosen – and a random action is chosen when no state is better than others. The controllers approach can be seen below:

```

for action : PossibleActions do
  | newGameState = gameState.copy().advance(action)
  | value[action] = value(newGameState)
end
return action leading to highest value, or random action if values are equal

```

Deep-search

This controller starts out in a similar fashion as the One-step, by expanding using all possible actions by copying the initial game-state. These game-states are then simulated further upon by advancing each game state a single time, with a random action, without copying the state (a rather costly procedure). This expansion is continued until the controller runs out of time, and a value for each action is calculated by considering the score and win/lose-value for each state that can be achieved from each of the initial states.

```
queue ← initial gameState
while has time left do
  gameState = q.poll()
  if gameState == initial gameState then
    for action : PossibleActions do
      | queue ← gameState.copy().advance(action)
    end
  else
    gameState.advance(random action)
    d ← depth of search (amount actions performed)
    value[action] += value(newGameState) * PossibleActions.lengthd
  end
end
return action leading to highest value, or random action if values are equal
```

Explorer

The explorer controller was designed specifically to play the arcade-style games of the GVG-AI framework, and utilizes a series of methods to strengthen its decisions. Unlike the other controllers which utilise open-loop searches, it stores information about visited tiles and prefers visiting unvisited locations. Overall the controller uses three different arrays to choose an action: One considering how probable death is from each action, one examining the score that can be achieved from the actions, and one only looking at the boringness. If the *death*-array have too high values, that array is used to take the decision (the action leading to the lowest value is chosen), otherwise the best action from the score-array is used. If the values of the *score*-array happen to be too similar, the *boringness*-array is lastly used instead.

The controller also addresses a common element of the VGDL example games, randomness. The controller gains an advantage in many of the games by simulating the results of actions repeatedly, before deciding the best move.

The Explorer was proven to be of a decent quality by getting a 2nd / 7th place in the GVG-AI competition. [SOURCE](#)

3.2.2 Puzzle controllers

Breadth first

Best first

3.3 FastVGDL

Since the GVG-AI framework has some functions and properties which are not interesting for some parts of this work, we created an implementation of a lighter version of VGDL, solely focused on analysing puzzle games in more detail. The implementation is basically a clone of the GVG-AI framework, but with several time- and memory consuming features removed, which in part was possible due to the puzzle games being relatively more simple (for instance, only movement from one tile to another is possible in FastVGDL, whereas sprites can move and collide in-between tiles in the GVG-AI framework).

SMALL TEST SHOWING TIME AND MEMORY DIFFERENCE PLOX!!!

Chapter 4

Automatic generation of VGDL descriptions

This chapter describes the setup- and methodology used in generating new games (i.e. game descriptions). Additionally an initial play test of set of both human-designed and generated games, using an array of different AI-controllers explained in the previous chapters.

This chapter is only focused on *action-arcade* games - *puzzle games* will be discussed and analysed in Chapter 6.

Selected designed arcade-action games

A subset of the human-designed games mentioned in Section 2.3 and 3.1 were selected for the following analyses games, and for creating new games by mutation.

After our addition of new games to GVG-AI framework we have access to 34 games, with five levels each. However, only 23 of the games are of the *arcade-action* genre explored in this section. Additionally, as mentioned in Section 2.3.1 a few of the example games from the GVG-AI competition are original creations, which cannot equally be assumed to be human-enjoyable.

With these limitation, thirteen interpretations of existing games are left and used as a baseline for testing and game generation: Aliens, Boulderdash, Frogs, Missile Command, Zelda, DigDug, Pacman, Seaquest, Eggomania, Solar Fox, Crackpots, Astrosmash and Centipede.

A slight change was also made to the GVG-AI framework, removing the possibility of players scoring lower than 0, which is the norm in the original games mentioned (this also makes comparing scores more straightforward).

4.1 Generating processes

Using the GVG-AI framework described in Section 2.3 we constructed a setup for generating large sets of new game description. We used two different approaches in generating new game description in VGDL: Mutating existing (designed) games, and randomly generating games from scratch. In both approaches several constraints were set upon the generation process to prevent crashes, and to increase the chance of a human playable game.

4.1.1 Mutation of example games

A process was built for parsing existing VGDL game descriptions, mutating certain sprites- or rules, and returning new, valid game descriptions,. The set of games described at the beginning of this chapter, were chosen to be used as a basis for mutating new game descriptions, analysed in this section.

Generating approach

It is not immediately apparent which approach to use in mutating the different elements of each VGDL game description. Since sprite- and rules are constructed from a combination of a class and

a series of parameters specific to the class, it is necessary to change the parameters if the class is changed, but the parameters of an existing class can freely be mutated.

A process was built for mutating each of the different parts of game descriptions, i.e. changing the sprites available of the `SpriteSet`, the interactions rules (`InteractionSet`) and the termination rules (`TerminationSet`). The process additionally allowed for changing the amount of rules, i.e. generating new rules or removing existing ones.

Since the sprites and rules of each game description is described by both a class and a series of parameters, several partly subjective decisions was made on defining probabilities of different parameters' values to try reach "realistic" values. For instance it was decided that a sprite being mutated or generated only has a 25% chance of using the `cooldown` parameter (making sprites have pauses between acting), and that the parameter have a random value between 1 and 10. The values and probabilities used were mostly decided by examining the set of designed games descriptions. Several constraints were also used to avoid game with non-valid descriptions (which can cause crashes in the GVG-AI framework), for instance by ensuring that avatar-sprites cannot be spawned and sprites cannot transform to an avatar, but an avatar sprite can still transform to another type of avatar sprite.

In this work we decide to only mutate interaction-rules from the `InteractionSet` of designed VGDL games, and simply change every part (classes, parameters and references) of a random amount of rules, with each rule having a 25% change of being changed for each game to mutate. In addition to simplifying the mutation process, and having games more with a higher chance of being playable this allows us to use the original designed games' levels for testing.

EXAMPLE OF MUTATED GAME — GAME DESCRIPTION AND/OR SCREENSHOT

4.1.2 Random game generation

A similar process as mentioned above was used to randomly put game descriptions together creating new games completely from scratch, constructing the textual lines for different parts of a VGDL description: Generating an array of sprites (for the `SpriteSet`), interaction-rules (`InteractionSet`), termination-rules (`TerminationSet`) and level mappings (`LevelMapping`).

Several partly subjective choices were again made for the range for the amount of different sprites, and interaction- and termination rules. Even though some of the original games contain a larger set of elements, we constricted the ranges to sprites: 3 – 8, interactions: 3 – 10 and terminations: 2 (a "win-" and a "lose"-termination).

When generating descriptions, we used similar constraints to those mentioned in section ??, partly to avoid generating descriptions with invalid elements, and partly to increase the proportion of interesting outcomes. As for the mutated games, some partly subjective decisions were made on defining the possible amount of sprites and rules, the proportion of sprites that has a level mapping, and when generating new sprite definitions.

To simplify the sprite creation process, no parent-child structure was used in the `SpriteSet`. However, the goal of the parent-child structure is only to make rule definitions less verbose, and does not make actual new game features possible.

All sprites were given a random sprite image, while the avatar-sprite was given the same sprite for each game, to make the game more easy to understand when visualizing the games (and actually play through them).

Level generation

As mentioned in Section 2.2.1 the problem of generating a game is more often than not intimately linked to a generation of levels. I.e. we could end up with generating game descriptions of high quality games, but if the level descriptions does not fit to the game-play (by being too trivial, or too hard for instance) the games might not found if searching through a set of randomly generated games.

A simple level generator were constructed with the simple goal of making the randomly generated games playable. The generator designates a level built using a given size, and inserts sprite mappings in different positions which takes all the level mappings from a game description and put a random (but small) amount of each sprite defined, in random location(s).

4.2 Initial test: Experimental setup

In this section we present a set of initial results from playing through both designed and generated games using different AI controllers, and discuss different ways to analyse the resulting data.

Controllers

Eight general videogame controllers were used to test the games. The controllers use different approaches, with a varying degree of intelligence. Four of the controllers are included in the GVG-AI framework, while the remaining were implemented for this work. Except for *OneStep-Heuristic*, the controllers only evaluate a given state according to its score and win/loss status. Each controller was marked with one of four types of intelligence for later tests: Intelligent (complex searches in game space are taken), semi-intelligent (game space search uses simple approach), random and do-nothing.

Generated games

A set of XXX mutated games were generated, by mutating each of the designed games ten times using the approach described in the previous section. 400 randomly generated games, each with a single generated level attached, were additionally developed and used in testing.

Testing and result analysis

The eight controllers were used to play through different set of example-, mutated and randomly generated games. Because of CPU budget limitations, each game was played with a maximum amount clock ticks of 800, and each controller was restricted to use 50 ms on each tick. For each playthrough only the default data from the GVG-AI framework was retrieved: The *score*, the *win-lose value*, the amount of *clock ticks* used and a list of actions performed.

A process to analyse the data was built, to calculate averages, standard deviations, max-, minimum and other statistics of each of the values, for each controller, with the possibility of calculating averages across several levels or games. To more accurately compare the *score* for the different controllers when playing across a range of different games we additionally calculate a normalised score using a max-min normalisation. The entropy of actions performed was also calculated, again with an average over all play throughs- or games.

Each game was played through 25 times, all within a single level attached to the game.

4.3 Initial test: Results

In this section we show results of the tests explained above, analyse the average of all play-throughs for each controller, and compare the results with each other. Before analysing and presenting the data, for all of the below tests we first remove games from the set which we consider either too difficult to analyse, or too simple too possibly be of interest – namely we remove games by three different considerations: 1) When a controller was disqualified (too many sprites- and/or interactions happening), 2) where the standard deviation of score was zero for all controllers, and 3) where the game always end in less than 50 clock ticks (which is the maximum lifetime of generated *Flicker*-sprites).

Designed games

The eight controllers were set to play through each of the human-designed VGDL games, with the constraints mentioned above.

We present the averages across each single game (can be seen in Appendix C) and averages across all of the games, seen in Figure 4.1. The distributions show that more intelligent controllers tend to have more success, with both a higher win-rate and a significantly higher average score. The *clock-ticks* and *action entropy* does not show any similar distribution for the collection of all games. Examining the results from each game it is noticeable that the profiles for these values (*clock-ticks* and *entropy*) are very different from game to game.

It should be noted that the uncertainties for average score, clock-ticks and action entropy assume a normal distribution of the data which in general is not true, and so these values should be taken with a grain of salt.

<i>controller</i>	<i>ave. score</i>	<i>norm. score</i>	<i>win-rate</i>	<i>ave ticks</i>	<i>ave. act entr.</i>
Explorer	42.94 (5.73)	0.80 (0.02)	0.45 (0.02)	514.66 (12.40)	0.89 (0.01)
MCTS	23.14 (4.07)	0.49 (0.02)	0.25 (0.02)	629.58 (11.30)	0.90 (0.01)
GA	11.98 (2.41)	0.40 (0.02)	0.15 (0.02)	590.34 (12.93)	0.71 (0.02)
Onestep-S	14.48 (1.29)	0.41 (0.02)	0.08 (0.01)	477.72 (15.09)	0.88 (0.02)
Onestep-H	3.73 (0.46)	0.24 (0.01)	0.11 (0.01)	569.27 (14.91)	0.19 (0.02)
Random	7.52 (0.80)	0.25 (0.01)	0.06 (0.01)	395.88 (14.97)	0.90 (0.01)
Do Nothing	0.39 (0.19)	0.13 (0.01)	0.06 (0.01)	601.72 (14.07)	0 (0)

Figure 4.1: Averaged results across all the XX designed games

Mutated games

After removing games by the requirements mentioned above, we were left with a total of 146 games in which we had useful data. The results of playing through these games with the controllers can be seen in Figure 4.2. The scores have higher means and standard deviations, indicating outliers in the data. The ordering of the *normalised score mean* and *win-rate*, however, shows a similar pattern as for the example games, with Explorer again excelling.

It should be noted here that the differences between the extreme values of *score* and *win-rate* are smaller than for the designed set of games.

<i>controller</i>	<i>ave. score</i>	<i>norm. score</i>	<i>win-rate</i>	<i>ave ticks</i>	<i>ave. act entr.</i>
Explorer	392.08 (106.62)	0.84 (0.01)	0.31 (0.01)	607.87 (4.20)	0.84 (0.01)
MCTS	140.49 (18.17)	0.48 (0.01)	0.15 (0.01)	689.05 (3.40)	0.90 (0.00)
GA	88.96 (12.53)	0.43 (0.01)	0.13 (0.01)	671.44 (3.83)	0.67 (0.01)
Onestep-S	128.82 (11.70)	0.44 (0.01)	0.10 (0.01)	565.77 (4.79)	0.84 (0.01)
Onestep-H	82.67 (12.63)	0.26 (0.01)	0.19 (0.01)	568.66 (4.89)	0.20 (0.01)
Random	69.70 (11.04)	0.26 (0.00)	0.11 (0.01)	470.10 (5.19)	0.90 (0.00)
Do Nothing	81.55 (15.65)	0.18 (0.01)	0.10 (0.00)	642.16 (4.16)	0 (0)

Figure 4.2: Averaged results across the 146 mutated games

Generated games

Figure 4.3 shows results for the remaining 65 randomly generated games (of 400), with problematic games removed according to the same criteria as in the previous section.

First of all, the *score* have much more extreme values than in the previous games, with the minimum being 199,406.58, over 1500 times larger than the highest in the set of example games (i.e. 121.55, by Explorer). Clearly, only the *normalised mean* can be used to compare scores across the different controllers. The *normalised score means* and *win-rates* both have values that are more closely clustered together, than in the previous game sets. For this set of games the *action entropy* of the intelligent controllers has also fallen quite a bit (relative to the Random-controller), indicating that more of the games have simple solutions to win or increase the score

It should be noted that the description generation process can often end up making games be automatically won (for instance, the goal of a game could be to destroy all coins, but all coins are of the **Flicker**-class and disappear after a certain amount of time), which causes the relative high *win-rate* for the Do Nothing-controller.

<i>controller</i>	<i>ave. score</i>	<i>norm. score</i>	<i>win-rate</i>	<i>ave ticks</i>	<i>ave. act entr.</i>
Explorer	−18 207.91 (5638.19)	0.62 (0.01)	0.26 (0.01)	592.43 (8.42)	0.72 (0.01)
MCTS	−4035.85 (6422.29)	0.44 (0.01)	0.28 (0.01)	590.21 (8.50)	0.84 (0.01)
GA	−3501.65 (6512.05)	0.44 (0.01)	0.25 (0.01)	607.55 (8.23)	0.58 (0.01)
Onestep-S	−16 680.67 (5745.31)	0.39 (0.01)	0.22 (0.01)	632.94 (7.74)	0.82 (0.01)
Onestep-H	−25 728.97 (4846.43)	0.36 (0.01)	0.20 (0.01)	644.59 (7.63)	0.40 (0.01)
Random	−23 348.24 (4946.64)	0.32 (0.01)	0.21 (0.01)	637.22 (7.69)	0.86 (0.01)
Do Nothing	−3051.34 (6438.95)	0.37 (0.01)	0.18 (0.01)	643.60 (7.70)	0 (0)

Figure 4.3: Averaged results across the 65 generated games

Outcome and graphs

By examining the graphs (Figure 4.4 and Figure 4.7) of the data just discussed, it is clear that there is a strong relationship between the performance profiles of playing with controllers of different types, on different set of games. However it should be noted here that some games from both the mutated- and generated game-set, contains games, which when examined by themselves have similar distributions as the average designed game.




Figure 4.4: Averaged normalised score across all games of the three different set of games: Designed-, mutated- and completely generated games

4.4 Discussion of initial test results

The results display some interesting patterns. Win rates suggest a relationship between intelligent controllers' success and better game design; for better designed games, the relative performance of different types of algorithms differ more. This corroborates our hypothesis that relative algorithm performance profiles can be used to differentiate between games of different quality. In randomly generated games, which arguably tend to be less interesting than the others, smarter controllers (e.g. Explorer and MCTS) do only slightly better than the worse ones (i.e. Random and DoNothing). This is due to a general lack of consistency between rules generated in this manner. Mutated games, however, derive from a designed game. Therefore, they maintain some characteristics of the original idea, which can improve the VGDL description's gameplay and playability.




Figure 4.5: Averaged win-rate for the three sets




Figure 4.6: Averaged clock tick for the three sets




Figure 4.7: Averaged action entropy for the three sets

While it is possible that random actions can result in good outcomes, this chance is very low, especially when compared to the chance of making informed decisions. In spite of that both Random and DoNothing do fairly well in randomly generated games. The performance of DoNothing emerges as a secondary indicator of (good) design: in human-designed games, DoNothing very rarely wins or even scores.

Human play of generated games

As mentioned in Section ?? there was a series of generated games with performance profile basically indistinguishable from the designed games. We decided on examining many of these games further by studying their game descriptions, and playing the games with visuals enabled, both with AI controllers and human players. (SCREEEEEEENSHTOS PLOX!!!) Besides a few mutated

games which kept most of the original rules (and therefore were successful games), the tested generated games with well-formed performance profiles were in general too trivial and aimless to be entertaining. This however gives an indication of an important assumption: That we can assume that the generated games, at least from the randomly generated set, are of a overall low quality. We will carefully use this assumption to construct an evaluation function for games, in the next chapter,

Besides the above, a few re-occurring problems was found in the generated games.

- Sprites often leave the level playing field.
- Several of the sprites- and/or rules are never used.
- The game can only be won in the first few (< 50) frames.
- Too much random stuff going on.

Chapter 5

Evaluation of games (fitness functions)

The initial tests of the previous chapter shows that there is reason to explore the relationship of AI's performance profiles to analyze a given VGD game, and give some indications of interesting statistical values to analyse.

This chapter focuses on, using the results of the previous chapter, creating a set of fitness functions to analyse different generated games, allowing a genetic program to evolve well-designed games, by creating and evolving VGD game descriptions. The resulting generated games are then discussed in detail.

5.1 Fitness function features: Introduction

From the above tests we can see that it is interesting to write a fitness function for controllers results, to be able to find out if new generated games are of a high quality. To make this possible we need to choose some set of values from the performance profiles of controllers from playing the games, to calculate an actual fitness value. As mentioned in Section ?? the *score* and *win-rate* seems to have the most distinct distributions for the different type of games, making them interesting to use for an evaluation function. The *action entropy* might also be appropriate to use, but the difference in distributions for that value is not as unique.

However even when limiting the features to score and win-rate, there is still an enormous amount of possible ways to calculate features-values for the fitness function. For instance, we could have a feature for the difference in average score between every single controller but it would result in $\binom{7}{2} = 21$ features for a single statistic. About statistics, we could use a series of different values simply to compare the score: *minimum*- and *maximum* score, *standard deviation* of score, *median*, *quartiles* or any form of normalised score, and of course we could use combined values for the features, e.g. the *score* increase per *clock tick*.

Additionally it is not clear which approach to use in actually calculating each feature. For instance, if we want to compare the analyse the score difference of the Explorer- and Random controller achieved in a given game, we need to compare two values that can have both positive and negative values.

5.2 Fitness function features: Selection

In this section we describe the process in which the fitness features were selected. As indicated in the previous section, a series of partly subjective decisions were needed for choosing features for fitness functions. To make an as informed decision as possible we ran a series of test described below, using different features, automatically choosing different set of features, evolving weights for the features, and examining the results using data mining principles.

In the following we decided to use a *relative difference*-function to compare the statistics for two different controllers' results.

5.2.1 Analysis of feature choices

As mentioned in the previous sections, the relationship of score and win-rate between controllers seems to be obvious choices, to to adopt as fitness features. Also, there seems to be a consistent difference in the values between the Explorer-, and the OneStep, Random and Do Nothing-controllers. To measure the strength of different choices from we made a small comparison of different circumstances that could potentially be used, counting how many of the designed (and generated) games that fulfil the conditions.

Below is the results of trying a few potential choices:

<i>data type</i>	Designed		Mutated		Rnd. gen.	
	<i>count</i>	<i>%</i>	<i>count</i>	<i>%</i>	<i>count</i>	<i>%</i>
Explorer.Mean \geq Random.Mean	17 (of 18)	94.4	139 (of 146)	95.2	42 (of 56)	75.0
Explorer.Mean \geq Do Nothing.Mean	16 (of 18)	88.9	131 (of 146)	89.7	38 (of 56)	67.9
Explorer.Winrate \geq Random.Winrate	18 (of 18)	100.0	140 (of 146)	95.9	56 (of 56)	100.0
Explorer.Winrate \geq Do Nothing.Winrate	17 (of 18)	94.4	141 (of 146)	96.6	55 (of 56)	98.2
Explorer.Median \geq Random.Median	17 (of 18)	94.4	138 (of 146)	94.5	45 (of 56)	80.4
Explorer.Median \geq Do Nothing.Median	16 (of 18)	88.9	134 (of 146)	91.8	41 (of 56)	73.2

Figure 5.1: Simple analysis of possible fitness values

The features above (and other sets) were analysed by data mining techniques – i.e classification processes were used to see if it was possible to differentiate between the designed- and generated set of games.

5.2.2 Selection and evolution of weights

In choosing a set of evaluation features it is natural to value each with a weight, making the addition from each feature count in a balanced fashion.

For most of the set of chosen features (when many features were selected) there was no clear indication of which values were more important, and so we continuously used a CMA-ES approach to evolve a set of weights. For this task another problem arise however: Generating a fitness function to score a chosen set of weights. Several different approaches were tried for this task, but all relying on calculating the fitness for each of designed- and randomly generated games, and making the overall fitness for the designed game set higher than for the generated games – using the assumption that the games from the generated set are of a low quality (discussed in Section ??).

5.2.3 Outcome

In the end we chose to only use the Explorer, OneStep-, Random- and Do Nothing controller in calculating fitness.

We decided on a set of 19 features, where 15 are calculated from the relative difference between a pair of controllers, and 4 values, each with a unique calculation. The chosen features were:

Score mean 3 features comparing the Explorer, with each of the other controllers.

Score median 3 features comparing the contr

Score quartion 1 3 features comparing the Explorer, with each of the other controllers.

Score mean 3 features comparing the Explorer, with each of the other controllers.

Score mean 3 features comparing the Explorer, with each of the other controllers.

In the end we chose the following set of values to compare: *ave. score*, *min. score*, *max. score* and *win-rate*.

5.3 Evolution of games: Setup

A process for generating evolving new VGDL games were built using the fitness function found in the previous section. Two different approaches were used to create new games: 1) Evolving from an existing, human-designed game description, using human-designed levels, and 2) evolving from randomly generated games, using randomly generated levels.

From the project it should be obvious that the first approach, using existing games and levels, have a much better chance of both achieving a high fitness-value and be an enjoyable game. However, the second approach allows for magnitudes of more options and possibilities, since it is not restricted to any set of sprites- or levels.

The figure below shows the process which was used in both evolution approaches.

HERE SHOULD BE A PICTURE OF THE GENERATION PROCESS

Generate game -> if not player runs out of time -> if not any sprites out of bounds -> get results -> if no disq -> if no low time wins -> if no all squares equal

5.4 Evolution of games: Results

The resulting game kindda suck

HERE GOES A GRAPH OF HOW FITNESS INCREASES OVER TIME IN EVOLUTION

5.5 Discussion of results

The resulting games are unfortunately pretty bad.

It would be interesting to retrieve more data from playing through each game. For instance, by examining the proportion of sprites- and rules that have been in use in the game, the amount of sprites that has left the level field, or the amount of objects that have been killed or created.

5.6 HUUUUSK Designed action-arcade games: Generated levels

Results of what happens when generating levels for the example games.

Chapter 6

Puzzle generation

6.1 Turn-based puzzle games

Definition of games

The games described in this section are games where the puzzles are the only game-play. There is no fast movement, or quick reaction time required to win. Additionally, because the games are described using the GVG-AI framework, the games focus around a player avatar which can only move up, down, left or right (also, in the games described below, all movement are from one game-tile to another).

6.1.1 Games

From the GVG-AI framework (Section 2.3) and my own additions (Section 3.1), there is a total of 10 puzzle games to use in the following analysis.

However, several of the games were deemed unfit for a complete examination, because either the levels were too difficult (Bolo Adventures and Chip's Challenge), that the game had features difficult to solve for a the puzzle-solving AI controllers (Bombuzal), or that the gameplay was too simple (Painter). The games used in the following are Bait, Brainman, The Citadel, Zen Puzzle, Sokoban, Modality

6.1.2 Puzzle solving AIs

As mentioned in (WHERE ITS MENTIONED), the controllers from the GVG-AI competition are not well-suited for playing puzzle games, because the goal can often only be achieved by applying a specific series of moves, which the controllers are not very good at. Two general puzzle solving algorithms were implemented. The overall goal of the controllers were to analyze the game as much as possibl, rather than finding a solution fast, or using as low memory as possible They controllers use fact that wall-sprites in the different games always push back the player, and so the controllers do not try to move into a wall – this is achieved by storing the positions of every wall-sprite at the start of the game. Also, the controllers store each path it is had tried. For each path a game state is calculated and stored, by finding the position and type of each (non-wall) sprite in the game. A path is cut off if the calculated game state is the same as has appeared before.

Breadth-first search A breadth-first search algorithm was implemented by using a queue and letting each node expand to the adjacent tiles, using the "tricks" described above. Additionally the controller was given a low- and high-memory option: In the high memor

Best-first search

6.1.3 Level generation

A setup for creating levels for a given game were constructed using an evolutionary algorithm. The level generator described in section ?? was used to generate simple levels for games. In addition two extra options were added: Wall-sprites and ground-sprites. This reduces a lot of troubles since

the avatar would otherwise be able to escape the level in a lot of the games. Ground-sprites signify which sprite should appear on empty locations.

The fitness function used in evolving levels was found by letting the two "puzzle solving AIs" play through the level.

Mutation

Crossover

Two different levels were constructed into a new, by going over each tile on the level and with 50% chance take the sprite residing in the two different original levels.

The setup was as follows:

```
while not at end of this document do
  read current;
  if understand then
    go to next section;
    current section becomes this one;
  else
    go back to the beginning of current section;
  end
end
```

Algorithm 1: How to write algorithms

6.1.4 Level generation for generated games

The Title
Put here cool text like what is going on in the wauw

This is subsubsection with title "Level generation for generated games".

Chapter 7

Conclusion

A huge problem in evolving games with the process used in this project, is that intelligent AI controllers are used which spend a lot of time in playing the through each game – just playing through a single game ten times takes the Explorer controller up to $2000 \text{ ticks} * 40 \text{ ms} * 10 = 800$ seconds.

A more clever approach might be to first probe the games with some more simple controllers, or decreasing the allowed time for intelligent controllers, allowing the games to be played through more quickly, and decide if the game is worth spending time on for a more precise analysis (using intelligent controllers).

It might be interesting to analyse games using a wider array of controllers, or analysing by letting the same controller (or controllers) play with a different amount of time allowed per clock-tick, or even using controllers designed to "play bad" (trying to die, decrease score) could be interesting

Also, it could be interesting to retrieve more extrinsic data from playing through a game (e.g. the proportion of interaction-rules that was triggered).

Bibliography

- Activision, Inc. Crackpots, 1983.
- Cameron Browne. *Automatic generation and evaluation of recombination games*. PhD thesis, Queensland University of Technology, 2008.
- Michael Cook and Simon Colton. Ludus ex machina: Building a 3d game designer that competes alongside humans. In *Proceedings of the 5th International Conference on Computational Creativity*, 2014.
- Marc Ebner, John Levine, Simon M Lucas, Tom Schaul, Tommy Thompson, and Julian Togelius. Towards a video game description language. *Dagstuhl Follow-Ups*, 6, 2013.
- José María Font, Tobias Mahlmann, Daniel Manrique, and Julian Togelius. Towards the automatic generation of card games through grammar-guided genetic programming. In *FDG*, pages 360–363, 2013.
- Gearbox Software. Borderlands 2, September 2012. URL <http://www.borderlands2.com/>.
- Michael Genesereth, Nathaniel Love, and Barney Pell. General game playing: Overview of the aaai competition. *AI magazine*, 26(2):62, 2005.
- Mark Hendrikx, Sebastiaan Meijer, Joeri Van Der Velden, and Alexandru Iosup. Procedural content generation for games: A survey. *ACM Transactions on Multimedia Computing, Communications, and Applications (TOMCCAP)*, 9(1):1, 2013.
- John Levine, Clare Bates Congdon, Marc Ebner, Graham Kendall, Simon M Lucas, Risto Miikkulainen, Tom Schaul, and Tommy Thompson. General video game playing. *Dagstuhl Follow-Ups*, 6, 2013.
- Antonios Liapis, Georgios N. Yannakakis, and Julian Togelius. Computational game creativity. In *Proceedings of the 5th International Conference on Computational Creativity*, 2014.
- Mattel Electronics. Astrosmash, 1981.
- Mojang. Minecraft, November 2011. URL <http://www.minecraft.net/>.
- Mark J. Nelson and Michael Mateas. Towards automated game design. In *AI*IA 2007: Artificial Intelligence and Human-Oriented Computing*, pages 626–637. Springer, 2007. Lecture Notes in Computer Science 4733.
- Mark J. Nelson, Julian Togelius, Cameron Browne, and Michael Cook. The search-based approach. In *Procedural Content Generation in Games: A Textbook and an Overview of Current Research*. Springer, 2014a. URL <http://www.pcgbook.com>. (To appear.).
- Mark J. Nelson, Julian Togelius, Cameron Browne, and Michael Cook. Chapter 6: Rules and mechanics. In *Procedural Content Generation in Games: A Textbook and an Overview of Current Research*. Springer, 2014b. URL <http://www.pcgbook.com>. (To appear.).
- Tom Schaul. A video game description language for model-based or interactive learning. In *Computational Intelligence in Games (CIG), 2013 IEEE Conference on*, pages 1–8. IEEE, 2013.
- Adam M Smith and Michael Mateas. Variations forever: Flexibly generating rulesets from a sculptable design space of mini-games. In *Proceedings of the 2010 IEEE Symposium on Computational Intelligence and Games*, pages 273–280, 2010.

- Julian Togelius and Jürgen Schmidhuber. An experiment in automatic game design. In *Proceedings of the 2008 IEEE Symposium on Computational Intelligence and Games*, pages 111–118, 2008.
- Julian Togelius, Georgios N. Yannakakis, Kenneth O. Stanley, and Cameron Browne. Search-based procedural content generation: A taxonomy and survey, 2011.
- Julian Togelius, Mark J. Nelson, and Antonios Liapis. Characteristics of generatable games. In *Proceedings of the 5th Workshop on Procedural Content Generation in Games*, 2014.
- Michael Toy and Glenn Wichman. *Rogue*, 1980.
- Derek Yu. Spelunky, September 2009. URL <http://www.spelunkyworld.com/>.

Appendices

Appendix A

GVG-AI example games

Below is a short summary of each of games, describing the winning conditions, and how the player can increase his/her score:

Aliens VGDL interpretation of the classic arcade game *Space Invaders*. A large amount of aliens are spawned from the top of the screen. The player wins by shooting all the approaching aliens.

Goal Destroy all the incoming aliens, without being hit by them or their projectiles.

Scoring Destroy all the incoming aliens, without being hit by them or their projectiles.

Boulderdash VGDL interpretation of *Boulder Dash*. The avatar has to dig through a cave to collect diamonds while avoiding being smashed by falling rocks or killed by enemies.

Butterflies The avatar has to capture all butterflies before all the cocoons are opened. Cocoons open when a butterfly touches them.

Chase About chasing and killing fleeing goats. However, if a fleeing goat encounters the corpse of another, it get angry and start chasing the player instead.

Digdug VGDL interpretation of *Dig Dug*. Avatar collects gold coins and gems, digs his way through a cave and avoid or shoot boulders at enemies.

Eggomania VGDL interpretation of *Eggomania*. Avatar moves from left to right collecting eggs that fall from a chicken at the top of the screen, in order to use these eggs to shoot at the chicken, killing it.

Firecaster Goal is to reach the exit by burning wood that is on the way. Ammunition is required to set things on fire.

Firestorms Player must avoid flames from hell gates until he finds the exit of a maze.

Frogs VGDL interpretation of *Frogger*. Player is a frog that has to cross a road and a river, without getting killed.

Infection Objective is to infect all healthy animals. The player gets infected by touching a bug. Medics can cure infected animals.

Missile Command VGDL interpretation of the classic arcade game *Missile Command*. Player has to destroy falling missiles, before they reach their destinations. If the player can save at least one city, he wins.

Overload Player must get to the level after collecting coins, but cannot collect too many coins, as he will be too heavy to traverse the exit.

Pacman A VGDL interpretation of *Pac-Man*. Goal is to clear a maze full with power pills and pellets, and avoid or destroy ghosts.

Portals Objective is to get to a certain point using portals to go from one place to another, while at the same time avoiding lasers.

Seaquest VGDL interpretation of *Seaquest*. Avatar is a submarine that rescue divers and avoids sea animals that can kill it. The goal is simply to a high score.

Survive Zombies Player has to flee zombies until time runs out, and can collect honey to kill the zombies.

Whackamole VGDL implementation of the classic arcade game *Whac-a-Mole*. Must collect moles that appear from holes, and avoid a cat that mimics the moles.

Zelda VGDL interpretation of *Legend of Zelda*. Objective is to find a key in a maze and leave the level. Player also has a sword to defend himself against enemies.

Camelrace Player needs to get to endpoint to win.

Sokoban The objective is to move the boxes to holes, until all boxes disappear or the time runs out.

Appendix B

Games designed during thesis

B.1 Action arcade games

Crackpots [1983] VGDL implementation of the classic arcade game *Whac-a-Mole*. Must collect moles that appear from holes, and avoid a cat that mimics the moles.

Solar Fox VGDL interpretation of *Legend of Zelda*. Objective is to find a key in a maze and leave the level. Player also has a sword to defend himself against enemies.

Astrosmash [1981] The player controls a laser cannon at the bottom of the screen, with the goal of shooting down as many incoming meteors, bombs and other objects. Points are earned by destroying objects, but lost if the objects reach the ground. The game ends if the laser cannon is hit a few times by the incoming objects.

Centipede The objective is to move the boxes to holes, until all boxes disappear or the time runs out.

B.2 Puzzle games

Bait VGDL interpretation of the classic arcade game *Space Invaders*. A large amount of aliens are spawned from the top of the screen. The player wins by shooting all the approaching aliens.

Goal Destroy all the incoming aliens, without being hit by them or their projectiles.

Scoring Destroy all the incoming aliens, without being hit by them or their projectiles.

The Citadel VGDL interpretation of *Boulder Dash*. The avatar has to dig through a cave to collect diamonds while avoiding being smashed by falling rocks or killed by enemies.

Bombuzal The avatar has to capture all butterflies before all the cocoons are opened. Cocoons open when a butterfly touches them.

Chip's Challenge About chasing and killing fleeing goats. However, if a fleeing goat encounters the corpse of another, it get angry and start chasing the player instead.

Bolo Adventures VGDL interpretation of *Dig Dug*. Avatar collects gold coins and gems, digs his way through a cave and avoid or shoot boulders at enemies.

(Real) Sokoban VGDL interpretation of *Eggomania*. Avatar moves from left to right collecting eggs that fall from a chicken at the top of the screen, in order to use these eggs to shoot at the chicken, killing it.

Brainman Goal is to reach the exit by burning wood that is on the way. Ammunition is required to set things on fire.

Modality Player must avoid flames from hell gates until he finds the exit of a maze.

Painter VGDL interpretation of *Frogger*. Player is a frog that has to cross a road and a river, without getting killed.

Zen Puzzle Objective is to infect all healthy animals. The player gets infected by touching a bug. Medics can cure infected animals.

Appendix C

Designed games results

This chapter shows the individual data for each game, for the test discussed in Section ??.

Aliens

<i>controller</i>	<i>score mean</i>	<i>std.dev.</i>	<i>normalised-mean</i>	<i>winrate</i>	<i>act-entropy</i>
Explorer	82.64	1.91	1.0000	1.0000	0
MCTS	69.16	4.47	0.8392	1.0000	0
GA	70.08	4.53	0.8499	0.9600	0
Onestep-S	55.08	12.74	0.6710	0.1600	0
Random	45.32	17.24	0.5555	0.0800	0
Do Nothing	−1.00	0.00	0.0000	0.0000	0

Boulderdash

<i>controller</i>	<i>score mean</i>	<i>std.dev.</i>	<i>normalised-mean</i>	<i>winrate</i>	<i>act-entropy</i>
Explorer	82.64	1.91	1.0000	1.0000	0
MCTS	69.16	4.47	0.8392	1.0000	0
GA	70.08	4.53	0.8499	0.9600	0
Onestep-S	55.08	12.74	0.6710	0.1600	0
Random	45.32	17.24	0.5555	0.0800	0
Do Nothing	−1.00	0.00	0.0000	0.0000	0